

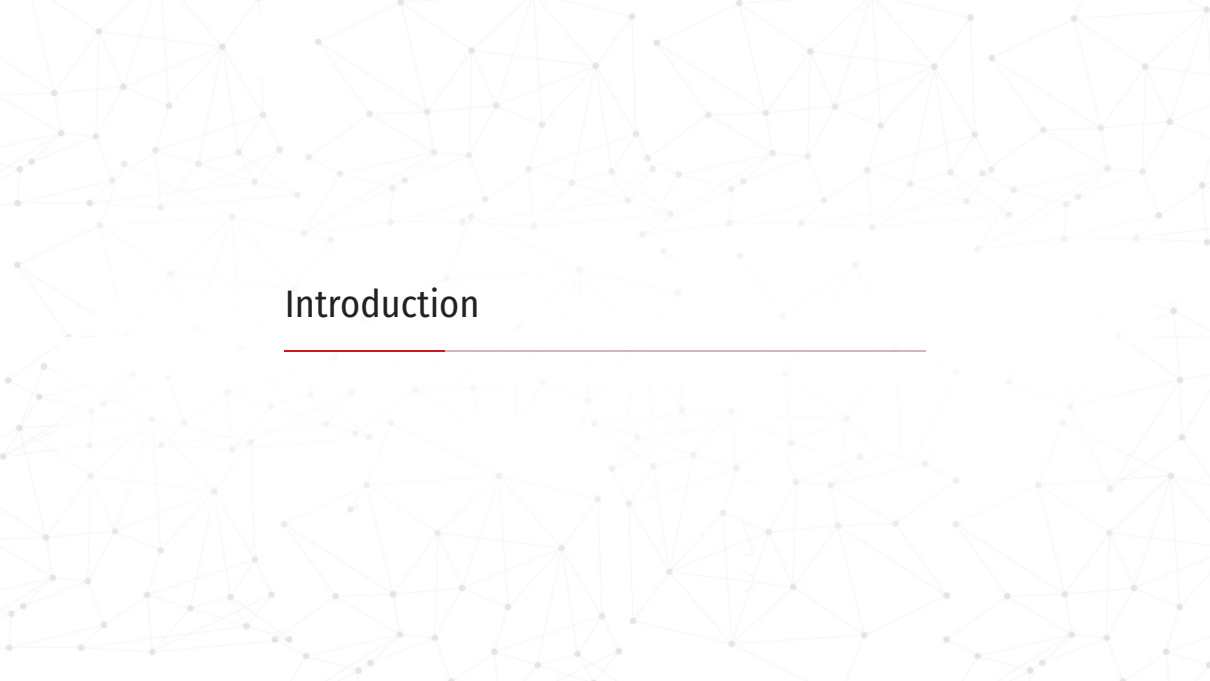


# Introduction to software-based microarchitectural side-channel attacks



Abc Xyz  
@dura\_lex

1. Introduction
2. Theory
3. Basic attacks
4. Software-based Microarchitectural Fault Attacks
5. Meltdown & Spectre
6. Summary

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These lines form a complex web of triangles and other polygons, creating a textured, mesh-like appearance across the entire slide.

# Introduction

---

Typical target of a side-channel attack



code1a:

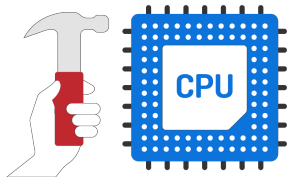
```
mov (X), %eax
```

```
mov (Y), %ebx
```


```
clflush (X)
```

```
clflush (Y)
```

```
jmp code1a
```



The DRAM cells get permanently damaged if hammered for a long time

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These connections form a complex web of triangles and other polygons, creating a textured, mesh-like appearance across the entire slide.

# Theory

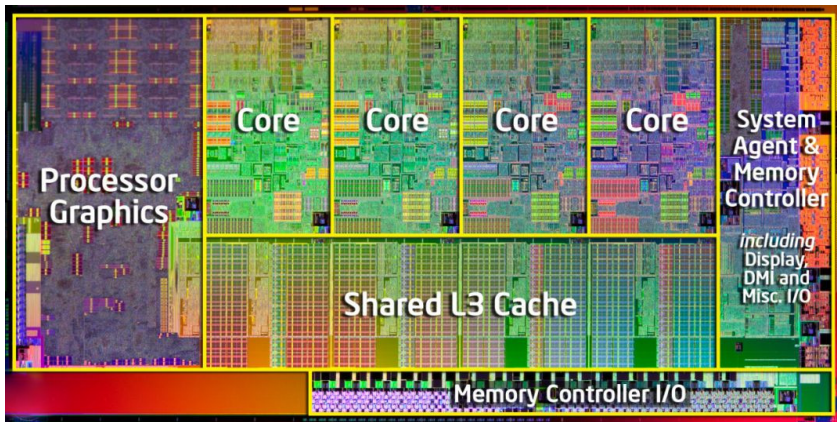
---

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These connections form a complex, web-like structure that fills the entire background, with some areas appearing denser than others.

Theory

---

CPU

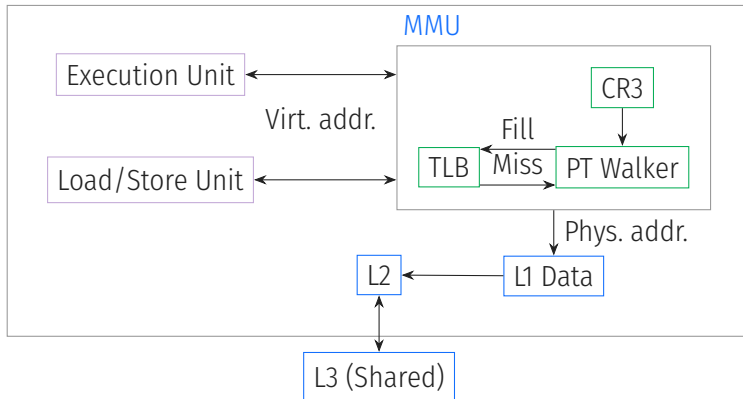


Architecture of multicore CPU



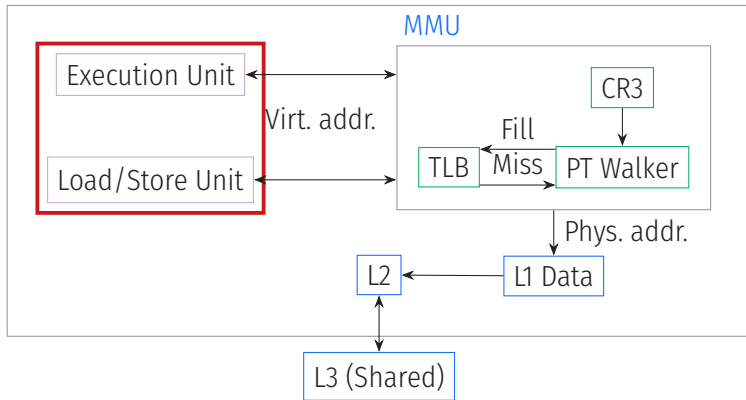
## Core

Abstract  
architecture of  
core and memory  
organization

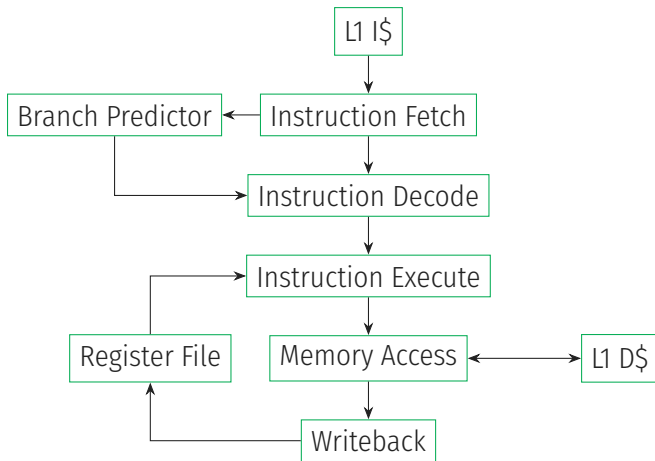


## Core

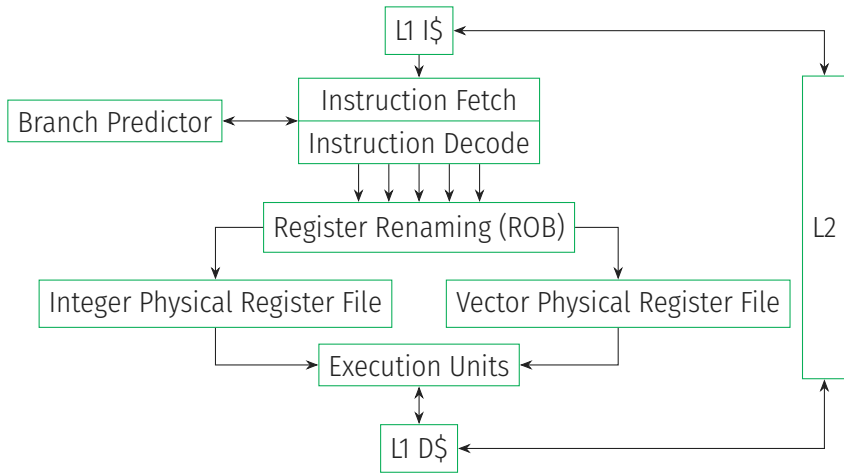
Abstract  
architecture of  
core and memory  
organization



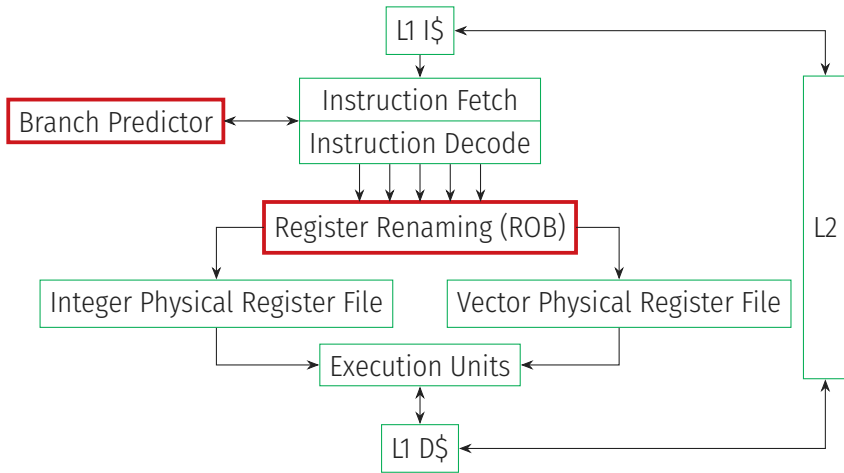
Elements of a modern  
in-order core

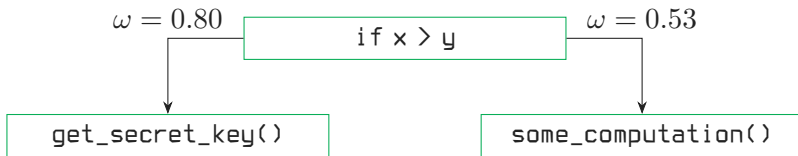


Elements of a  
modern  
out-of-order core

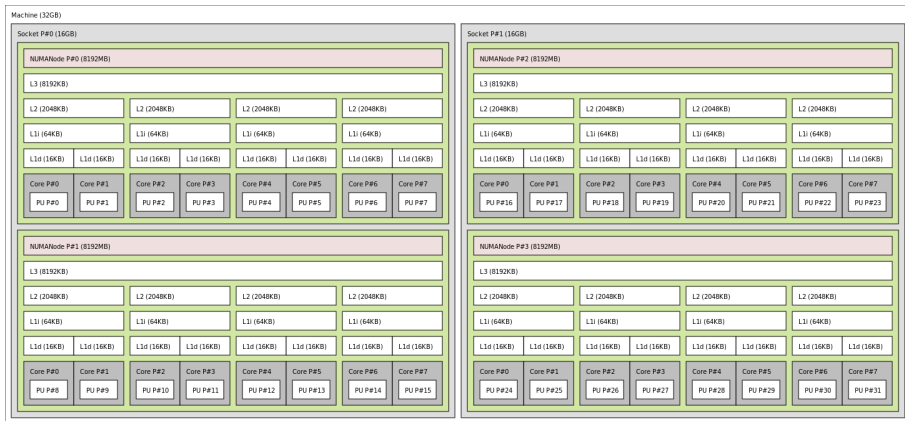


Elements of a  
modern  
out-of-order core





`get_secret_key()` can be executed speculatively



Architecture of multicore CPU AMD Bulldozer

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These connections form a complex, web-like structure that fills the entire background. The density of the connections is higher in some areas and lower in others, creating a sense of organic growth or a distributed system.

Theory

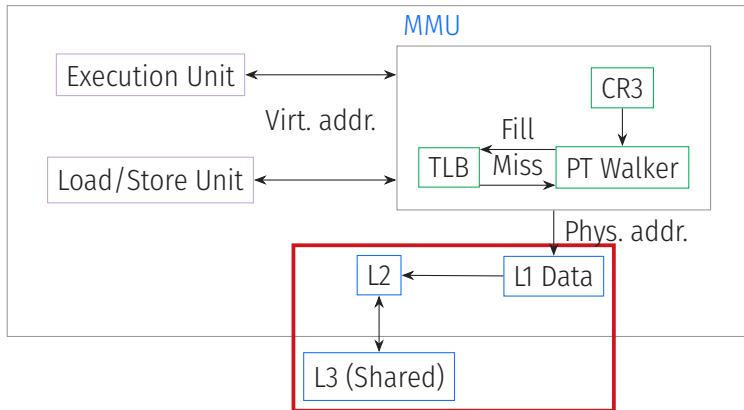
---

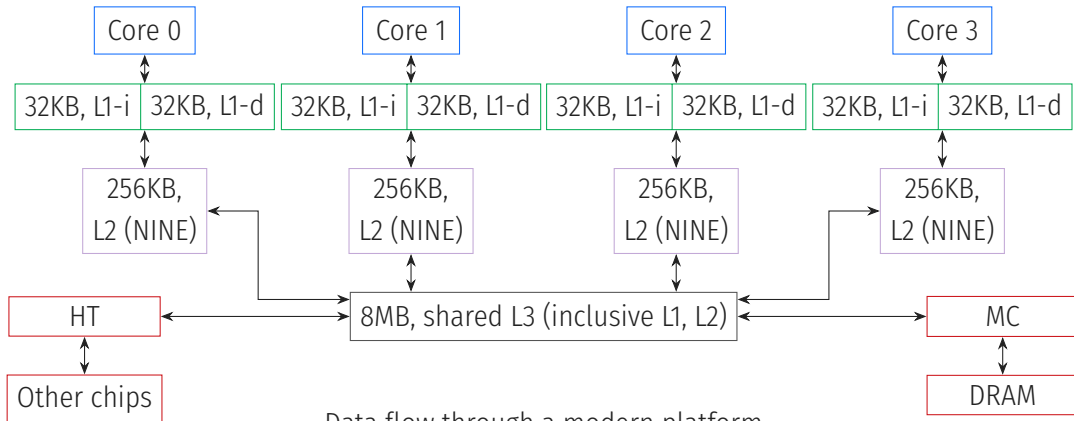
Cache



## Core

Abstract  
architecture of  
core and memory  
organization





Data flow through a modern platform

Virtual memory

0xf200
0xf100
0xf000
...
0x3000
0x2000

hit

hit

Address

Data

0xf200

Kernel secret 0

0x2000

User secret

Cache (L1, L2, LLC)

Physical memory (DRAM)

0x1000
0x0900
0x0800
0x0700
0x0600
0x0500

CPU cache algorithm

Virtual memory

0xf200
0xf100
0xf000
...
0x3000
0x2000

miss

Address	Data
0xf200	Kernel secret 0
0x2000	User secret

Cache (L1, L2, LLC)

Physical memory (DRAM)

0x1000
0x0900
0x0800
0x0700
0x0600
0x0500

CPU cache algorithm

Virtual memory

0xf200
0xf100
0xf000
...
0x3000
0x2000

Address

Data

0xf200	Kernel secret 0
0x2000	User secret

Cache (L1, L2, LLC)

Physical memory (DRAM)

0x1000
0x0900
0x0800
0x0700
0x0600
0x0500

CPU cache algorithm

Virtual memory

0xf200
0xf100
0xf000
...
0x3000
0x2000

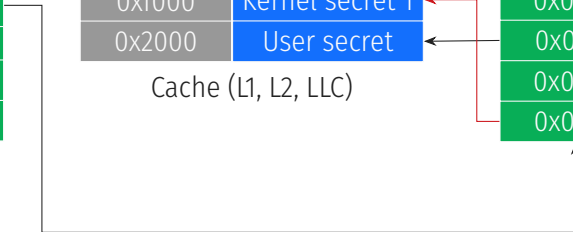
Physical memory (DRAM)

0x1000
0x0900
0x0800
0x0700
0x0600
0x0500

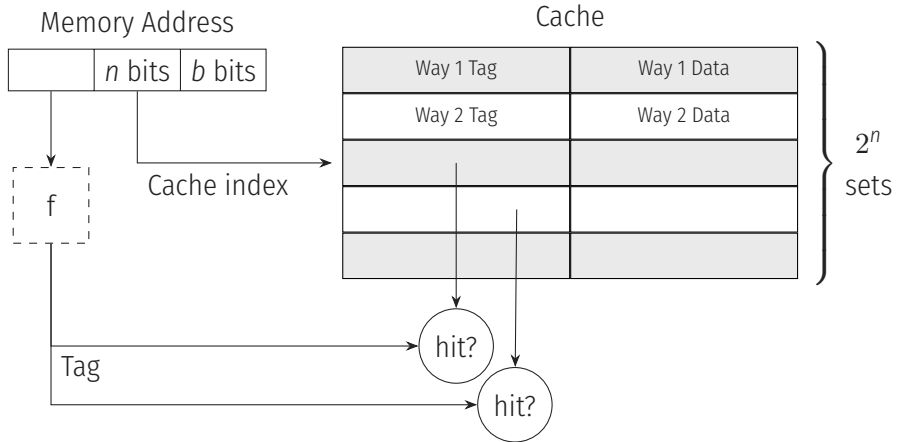
Address	Data
0xf000	Kernel secret 1
0x2000	User secret

Cache (L1, L2, LLC)

CPU cache algorithm




- Direct-mapped cache
- Fully-associative cache
- 2/4/8/12-way set associative cache





- FIFO
- LIFO
- least recently used, LRU
- time aware least recently used, TLRU
- most recently used, MRU
- pseudo-LRU, PLRU
- random replacement, RR
- segment LRU, SLRU
- least frequently used, LFU
- least frequent recently used, LFRU
- LFU with dynamic aging, LFUDA
- low inter-reference recency set, LIRS
- adaptive replacement cache, ARC
- clock with adaptive replacement, CAR
- multi queue, MQ
- and etc.

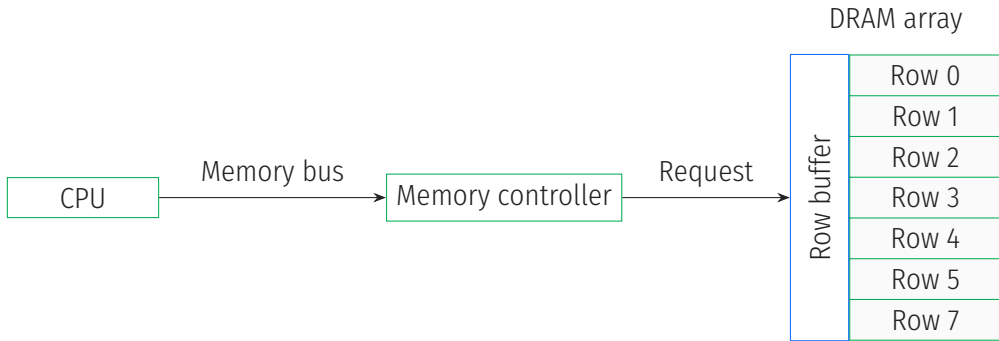
- Virtually indexed, virtually tagged (VIVT)
- Physically indexed, virtually tagged (PIVT)
- Virtually indexed, physically tagged (VIPT)
- Physically indexed, physically tagged (PIPT)

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These connections form a complex, web-like structure that fills the entire background, with some areas appearing denser than others.

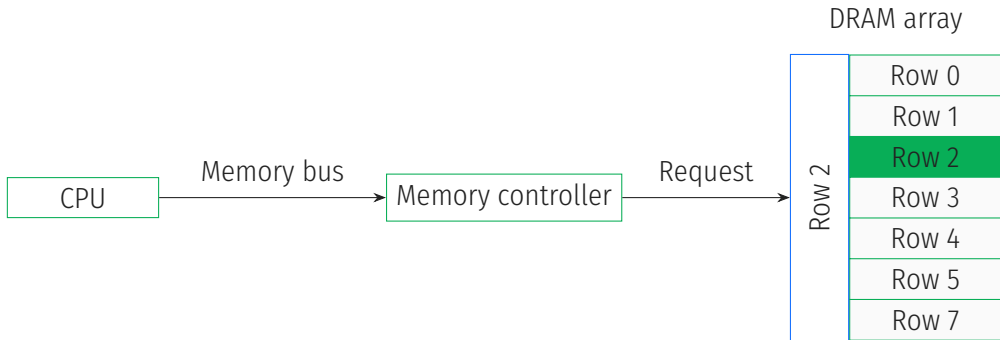
Theory

---

DRAM

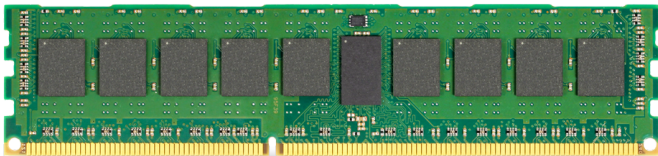


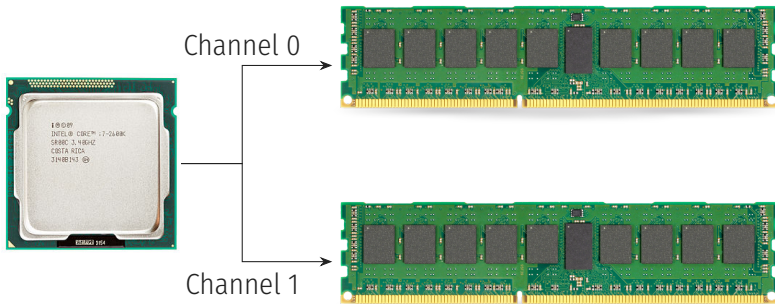
A simple computer system with a single DRAM array

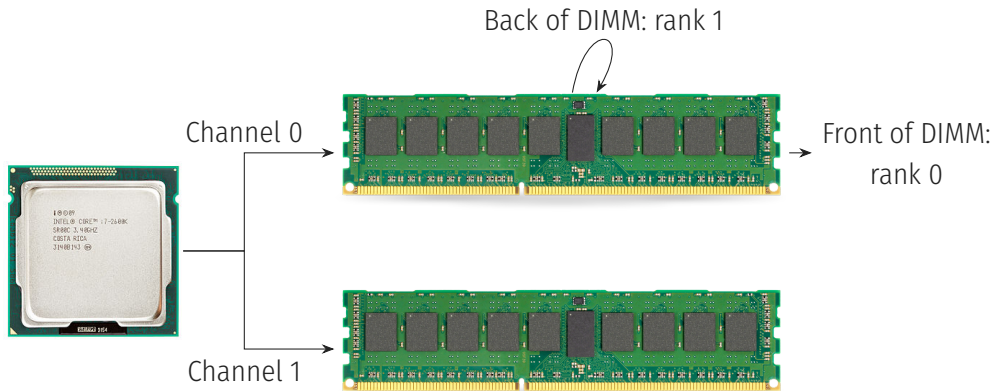


A simple computer system with a single DRAM array

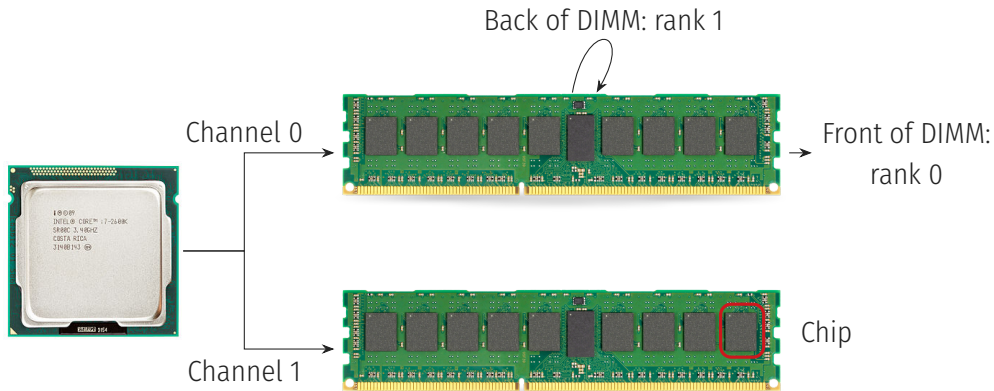
DIMM

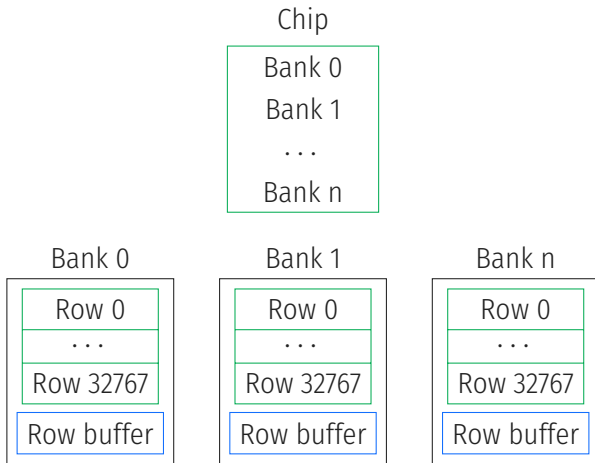


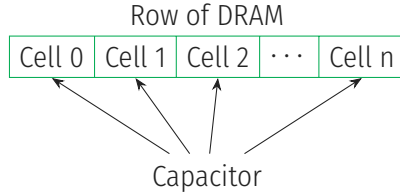












The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These connections form a complex, web-like structure that fills the entire background. The density of the connections is higher in some areas and lower in others, creating a sense of organic growth or a distributed system.

## Basic attacks

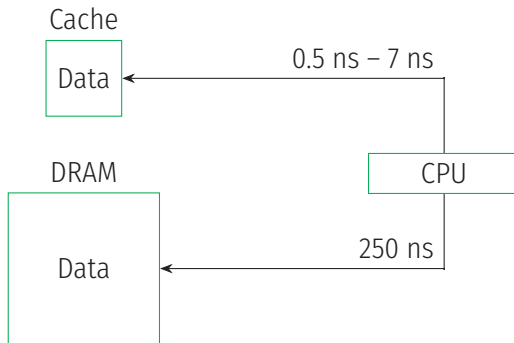
---

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines (edges), creating a web-like structure that covers the entire area.

## Basic attacks

---

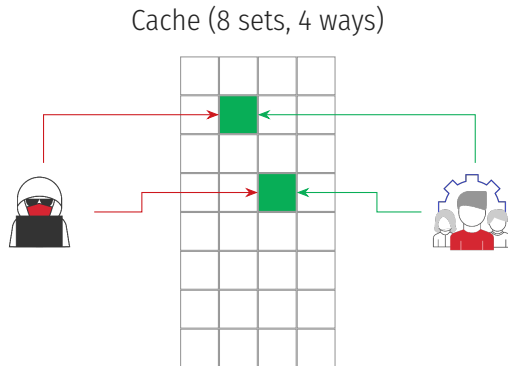
Cache attacks



Timing attack — an attack aimed at exploiting differences in an algorithm execution time

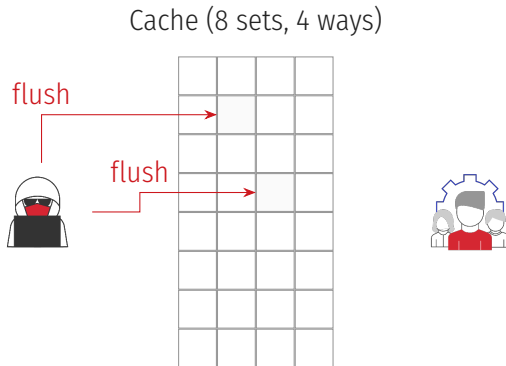
1. Map binary (e.g., shared object) into the address space
2. Flush a cache line (code or data) from the cache
3. Schedule a victim's program
4. Check if the corresponding cache line from Step 2 has been loaded by the victim's program

Map binary (e.g., shared object) into the address space





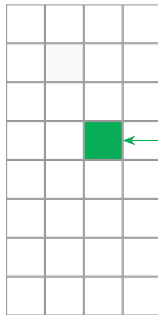
Flush a cache line (code or data)  
from the cache



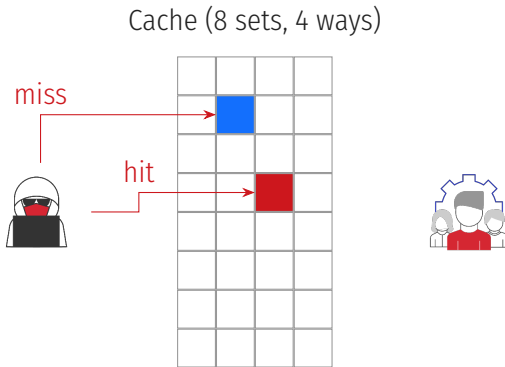
Schedule a victim's program



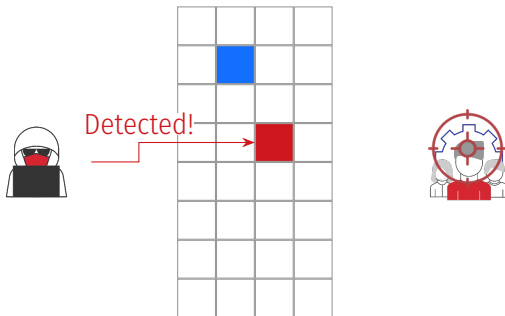
Cache (8 sets, 4 ways)




Check if the corresponding cache line from Step 2 has been loaded by the victim's program



Cache (8 sets, 4 ways)



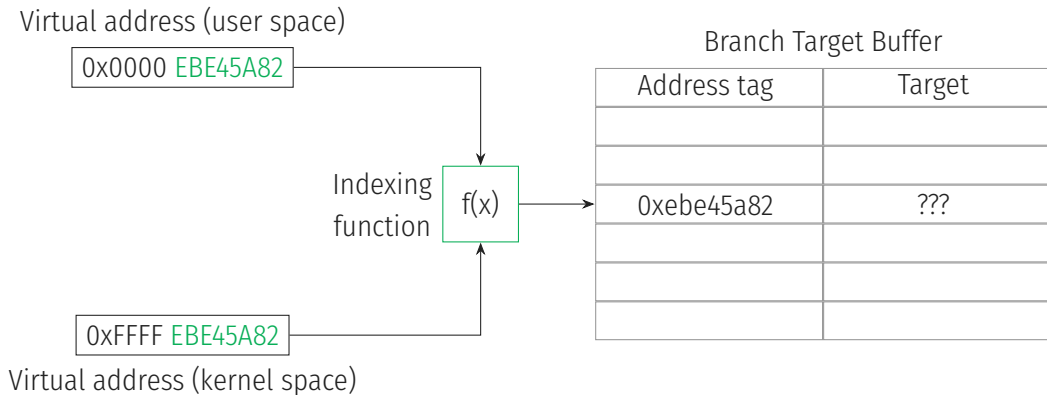
- Evict + Time
- Prime + Probe
- Prime + Abort
- Flush + Flush
- Evict + Reload
- AnC (ASLR  $\oplus$  Cache)
- etc.

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines (edges), creating a web-like structure that fills the entire frame. The density of the connections is higher in some areas and lower in others, giving it a dynamic, organic feel.

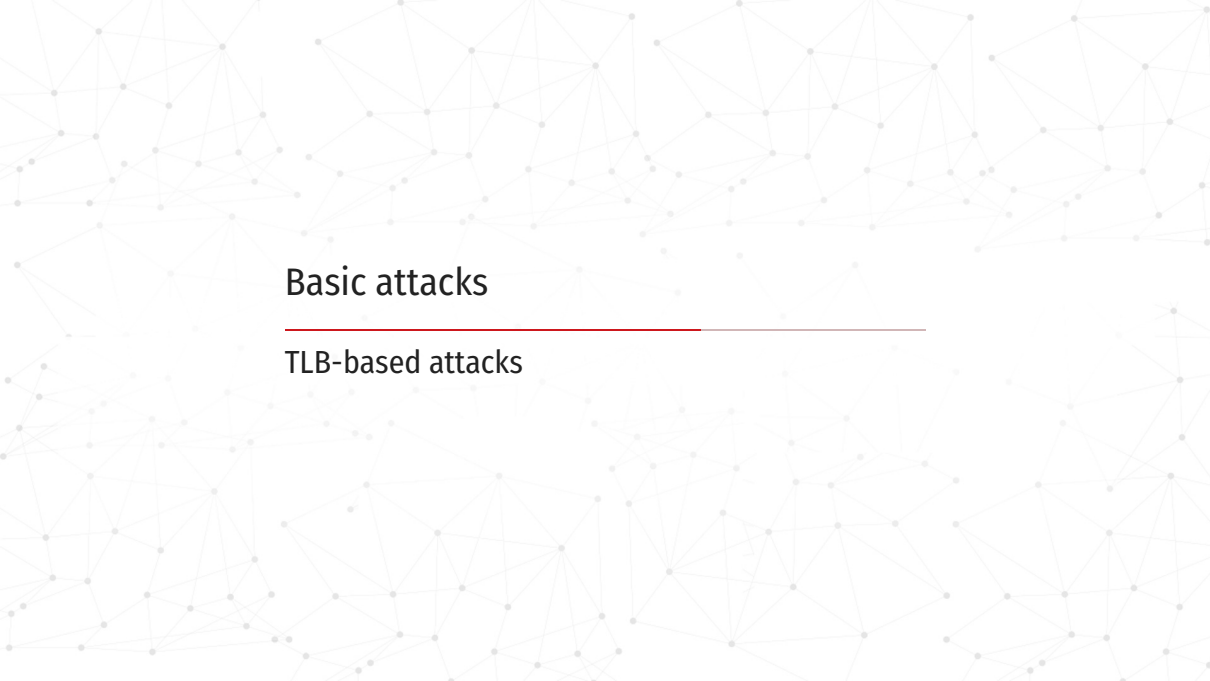
## Basic attacks

---

### Branch-prediction attacks



Branch Target Buffer addressing a scheme in the Haswell processor

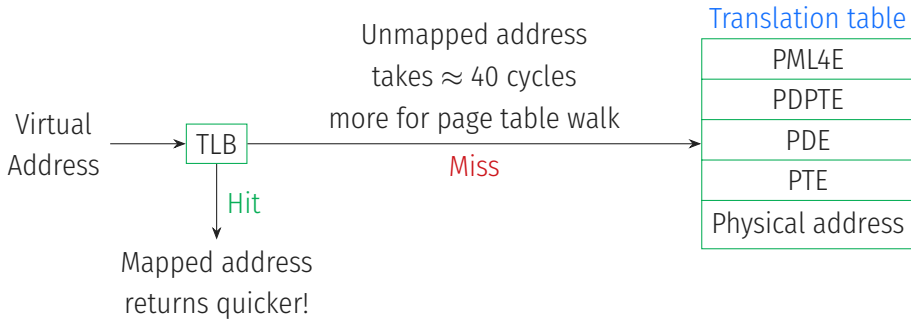
The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines (edges), creating a web-like structure that fills the entire frame. The pattern is more dense in some areas and sparser in others, giving it a dynamic, organic feel.

## Basic attacks


---

TLB-based attacks





A translation lookaside buffer (TLB) is a memory cache that is used to reduce the time taken to access a user memory location


The background of the slide features a complex, light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines, creating a web-like structure that fills the entire frame. The density of the connections varies, with some areas appearing more clustered than others.

## Basic attacks

---

Exception-based attacks

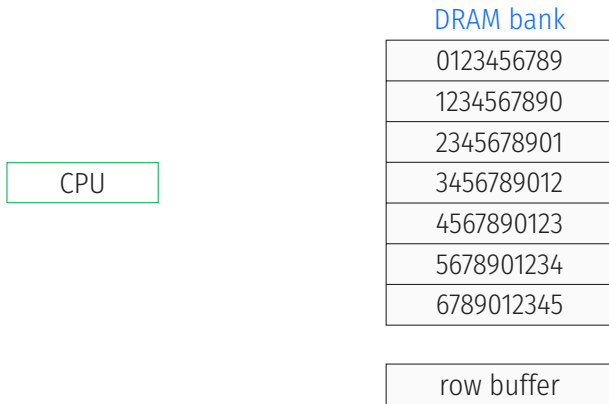
- Scheduler interrupts
- Instruction aborts
- Page faults
- Behavioral differences (e.g, error code)

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines (edges), creating a web-like structure that fills the entire frame. The pattern is more dense in some areas and sparser in others, giving it a dynamic, organic feel.

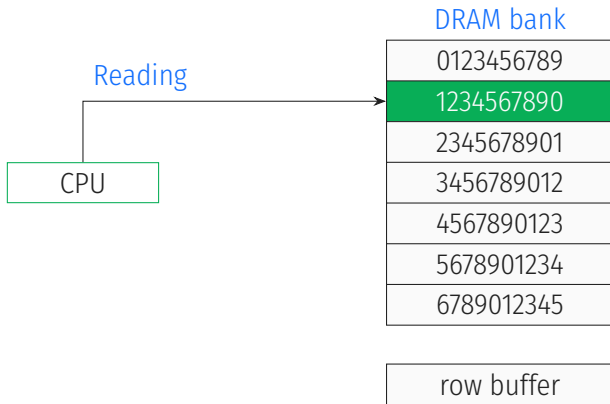
## Basic attacks

---

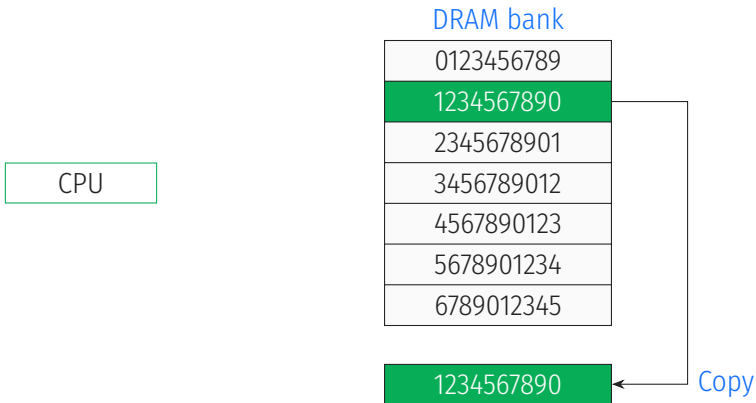
DRAM-based attacks

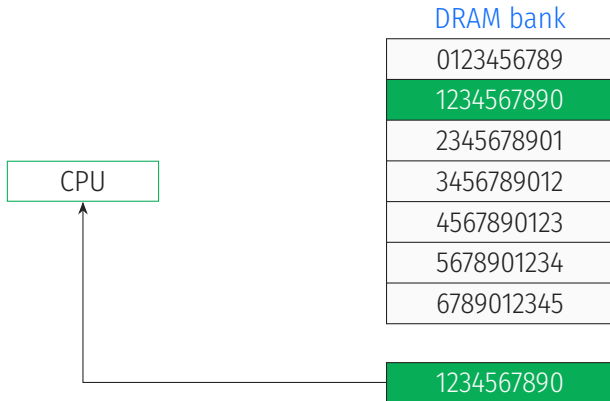


Reading from DRAM

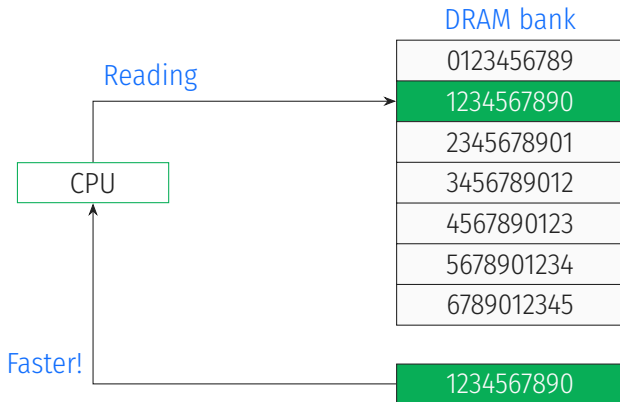


CPU reads row 1, row buffer is empty










CPU reads row 1, row buffer is now full

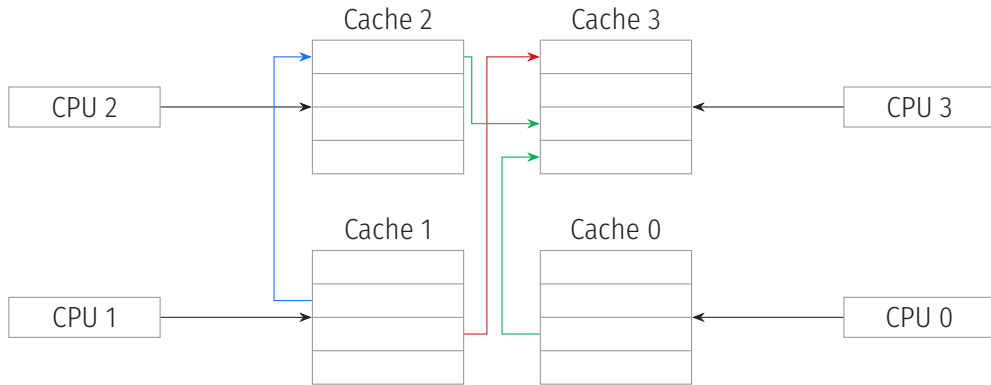
- DRAMA
- Row hit (Flush + Reload)
- Row miss (Prime + Probe)
- etc.

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines (edges), creating a web-like structure that fills the entire frame. The pattern is more dense in some areas and sparser in others, giving it a dynamic, interconnected appearance.

## Basic attacks


---

Covert channels



Cross-core covert channels

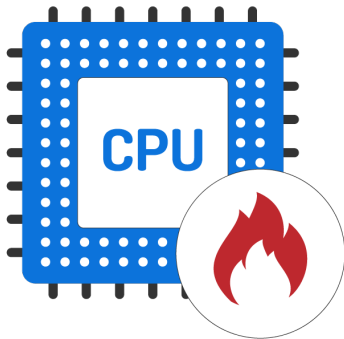
- Cache-based covert channels (shared libraries)
- Row miss attack (DRAM)
- Thermal covert channels
- Radio covert channels

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines, creating a web-like or molecular structure that covers the entire area.


# Software-based Microarchitectural Fault Attacks

---

# Software-based Microarchitectural Fault Attacks



Software-based microarchitectural fault attacks do not require physical access, but only some form of code execution on the target system

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines, creating a web-like or molecular structure that covers the entire area.

# Software-based Microarchitectural Fault Attacks

---

Rowhammer



- Fast uncached memory access
- Physical memory massaging
- Physical memory addressing

- Flip Feng Shui — targeted Rowhammer
- Throwhammer — remote Rowhammer
- Nethammer — better remote Rowhammer
- Drammer, RAMpage — exploitation ARM-based hardware
- Glitch — better exploitation ARM-based hardware

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These lines form a complex web of triangles and other polygons, creating a mesh-like texture across the entire slide.

# Meltdown & Spectre

---

The background of the slide features a complex, light gray network pattern. It consists of numerous small dots (nodes) connected by thin, intersecting lines, creating a web-like or molecular structure that covers the entire area.

# Meltdown & Spectre

---

Derived attacks and not only

- MeltdownPrime & SpectrePrime
- SgxPectre
- SMM Speculative Execution Attacks
- BranchScope
- LazyFP
- ...

- Spectre 1.1, 1.2 (Speculative Buffer Overflows)
- SpectreRSB
- NetSpectre
- L1TF (Foreshadow)
- etc.

TotalMeltdown? And what's about other patches...



# Meltdown & Spectre

---

Abstract example of exploitation

Four components of speculative techniques

1. Speculation primitive



## Four components of speculative techniques

### 1. Speculation primitive

- Bypass out of bounds checks
- Training of branch predictor
- Speculatively read an earlier value of the data
- Pending exceptions
- Exploit branch history table
- Exploit the Return Stack Buffer
- Speculatively write to register (buffer overflow)

Type of BP

Algorithm of BP

Environment of BP

The foundation of the Speculative-Based Attack tower

## Four components of speculative techniques

1. Speculation primitive
2. Windowing gadget

## Four components of speculative techniques

1. Speculation primitive
  2. Windowing gadget
- Non-cached loads
  - Dependency chain of loads
  - Dependency chain of integer ALU operations



The Speculative-Based Attack tower

## Four components of speculative techniques

1. Speculation primitive
2. Windowing gadget
3. Disclosure gadget

## Four components of speculative techniques

1. Speculation primitive
2. Windowing gadget
3. Disclosure gadget

- ASLR
- CFI
- SMAP
- DEP/NX
- retpoline
- and others.



The Speculative-Based Attack **Babel** tower



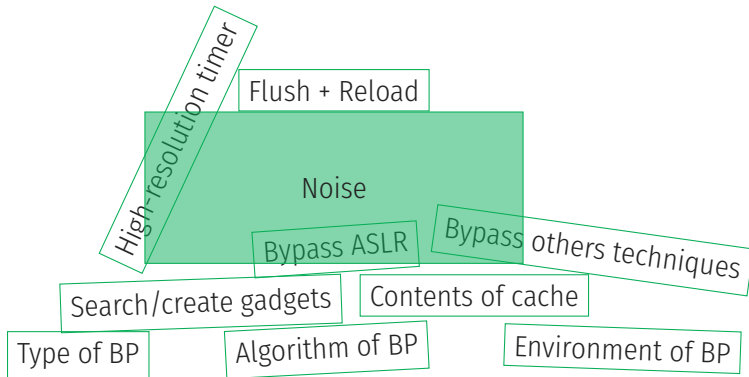
## Four components of speculative techniques

1. Speculation primitive
2. Windowing gadget
3. Disclosure gadget
4. Disclosure primitive


## Four components of speculative techniques

1. Speculation primitive
2. Windowing gadget
3. Disclosure gadget
4. Disclosure primitive

- Architecture of cache
- Replacement policies
- Exclusive and inclusive
- Type of cache attack
- Noise
- High-resolution timer
- and etc.



The Speculative-Based Attack **Babel** tower

The background of the slide is a light gray network pattern. It consists of numerous small, dark gray circular nodes connected by thin, light gray lines. These lines form a complex web of triangles and other polygons, creating a textured, mesh-like appearance across the entire slide.

## Summary

---

- Software-based microarchitectural attacks has become **very popular**

- Software-based microarchitectural attacks has become **very popular**
- Requires **a lot of resources** to develop a working exploit








- Software-based microarchitectural attacks has become **very popular**
- Requires **a lot of resources** to develop a working exploit
- Microarchitectural attacks may be **automated**






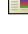
- Software-based microarchitectural attacks has become **very popular**
- Requires **a lot of resources** to develop a working exploit
- Microarchitectural attacks may be **automated**
- Many attacks have **not yet been disclosed**













- Software-based microarchitectural attacks has become **very popular**
- Requires **a lot of resources** to develop a working exploit
- Microarchitectural attacks may be **automated**
- Many attacks have **not yet been disclosed**
- Countermeasures come with a **performance impact**








Questions?







-  D. Gruss, “Software-based Microarchitectural Attacks.”
-  M. Lipp and D. Gruss, “ARMageddon: Cache Attacks on Mobile Devices.”
-  D. Page, “MASCAB: a Micro-Architectural Side-Channel Attack Bibliography.”
-  P. Pessl and D. Gruss, “DRAMA: Exploiting DRAM Addressing for Cross-CPU Attacks.”
-  H. Bos and Y. Fratantonio, “Drammer: Deterministic Rowhammer Attacks on Mobile Platforms.”
-  Microsoft, “Mitigating speculative execution side channel hardware vulnerabilities.”
-  J. Horn, “Reading privileged memory with a side-channel.”

-  D. Gruss and M. Lipp, “KASLR is Dead: Long Live KASLR.”
-  D. Gruss and C. Maurice, “Flush+Flush: A Fast and Stealthy Cache Attack.”
-  F. Liu and Y. Yarom, “Last-Level Cache Side-Channel Attacks are Practical.”
-  C. Trippel, D. Lustig, and M. Martonosi, “MeltdownPrime and SpectrePrime: Automatically-Synthesized Attacks Exploiting Invalidation-Based Coherence Protocols.”
-  M. Schwarz, C. Maurice, D. Gruss, and S. Mangard, “Fantastic Timers and Where to Find Them: High-Resolution Microarchitectural Attacks in JavaScript.”
-  M. Lipp and M. T. Aga, “Nethammer: Inducing Rowhammer Faults through Network Requests.”

-  A. Tatar and R. Krishnan, “Throwhammer: Rowhammer Attacks over the Network and Defenses.”
-  G. Camurati and S. Poeplau, “Screaming Channels: When Electromagnetic Side Channels Meet Radio Transceivers.”
-  J. Stecklina and T. Prescher, “LazyFP: Leaking FPU Register State using Microarchitectural Side-Channels.”
-  M. Guri and A. Kachlon, “GSMem: Data Exfiltration from Air-Gapped Computers over GSM Frequencies.”
-  D. Sullivan, O. Arias, T. Meade, and Y. Jin, “Microarchitectural Minefields: 4K-Aliasing Covert Channel and Multi-Tenant Detection in IaaS Clouds.”

-  B. Gras, K. Razavi, E. Bosman, H. Bos, and C. Giuffrida, “ASLR on the Line: Practical Cache Attacks on the MMU.”
-  S. van Schaik, C. Giuffrida, H. Bos, and K. Razavi, “Malicious Management Unit: Why Stopping Cache Attacks in Software is Harder Than You Think.”
-  D. Gruss and A. Fogh, “Prefetch Side-Channel Attacks: Bypassing SMAP and Kernel ASLR.”
-  E. M. Koruyeh and K. N. Khasawneh, “Spectre Returns! Speculation Attacks using the Return Stack Buffer.”
-  G. Maisuradze and C. Rossow, “ret2spec: Speculative Execution Using Return Stack Buffers.”

-  G. Chen and S. Chen, “SgxPectre Attacks: Leaking Enclave Secrets via Speculative Execution.”
-  M. Lipp and M. Schwarz, “Meltdown.”
-  P. Kocher and D. Genkin, “Spectre Attacks: Exploiting Speculative Execution.”
-  ARM, “Cache Speculation Side-channels.”
-  M. Schwarz, M. Schwarzl, M. Lipp, and D. Gruss, “NetSpectre: Read Arbitrary Memory over Network.”
-  S. D’Antoine, “Out-of-Order Execution and Its Applications.”
-  V. Kiriansky and C. Waldspurger, “Speculative Buffer Overflows: Attacks and Defenses.”

-  B. Gras, K. Razavi, H. Bos, and C. Giuffrida, “Translation Leak-aside Buffer: Defeating Cache Side-channel Protections with TLB Attacks.”
-  C. Disselkoen, D. Kohlbrenner, L. Porter, and D. Tullsen, “Prime+Abort: A Timer-Free High-Precision L3 Cache Attack using Intel TSX.”
-  M. Lipp and M. Schwarz, “Meltdown & Spectre Side-channels considered hARMful.”
-  J. Masters, “Exploiting modern microarchitectures: Meltdown, Spectre, and other attacks.”
-  M. Lipp, “Cache attacks on ARM.”
-  D. Evtvyushkin, D. Ponomarev, and N. Abu-Ghazaleh, “Jump over ASLR: attacking branch predictors to bypass ASLR.”